

Prediction of Sales Volumes

G. L.

15/10/2018

1. The Dataset: Products and Attributes

A company selling electronic items is interested in estimating the market potential of a series of new products that it plans to launch to the market. The sales potential of these new products can be predicted using information of similar products already in the market. This information is available as a dataset containing various types of products and their attributes. Specifically, the table contains:

- 245 electronic products of various types (computers, displays, consoles, etc.).
- 22 variables describing various product attributes.

```
wd <- file.path('~',
                'GitRepos',
                'r-ds-projects',
                'sales_predictions')
setwd(dir= wd)

library(readr)
library(stringr)
library(ggplot2)
library(ggpubr)
library(dplyr)
library(corrplot)
library(caret)
library(gbm)
library(doMC)

# parallelization
registerDoMC(cores = 4)

# read dataset of existing products
input_file <- file.path('.',
                        'existing_products.csv')
if( !file.exists(input_file) ) {
  print('File:')
  print(input_file)
  print('not found!. Current dir:')
  getwd()
}
prod_exist <- read.csv(input_file, dec= ',', sep= ';')
```

A look at the structure of the dataset.

```
str(prod_exist)

## 'data.frame':   245 obs. of  22 variables:
##  $ X                      : int  1 2 3 4 5 6 7 8 9 10 ...
##  $ Product_type           : Factor w/ 12 levels "Accessories",...: 7 7 7 5 5 1 1 1 1 1 ...
##  $ Product_ID             : int  101 102 103 104 105 106 107 108 109 110 ...
##  $ Prices                 : num  949 2250 399 410 1080 ...
```

```
## $ X5Stars          : int  3 2 3 49 58 83 11 33 16 10 ...
## $ X4Stars          : int  3 1 0 19 31 30 3 19 9 1 ...
## $ X3Stars          : int  2 0 0 8 11 10 0 12 2 1 ...
## $ X2Stars          : int  0 0 0 3 7 9 0 5 0 0 ...
## $ X1Stars          : int  0 0 0 9 36 40 1 9 2 0 ...
## $ Positive_service_review : int  2 1 1 7 7 12 3 5 2 2 ...
## $ Negative_service_review : int  0 0 0 8 20 5 0 3 1 0 ...
## $ Would_consumer_recomend__product: num  0.9 0.9 0.9 0.8 0.7 0.3 0.9 0.7 0.8 0.9 ...
## $ Best_seller_rank      : num  1967 4806 12076 109 268 ...
## $ Weighth              : num  25.8 50 17.4 5.7 7 1.6 7.3 12 1.8 0.75 ...
## $ Depth                : Factor w/ 138 levels "0","0.04","0.07",...: 97 108 47 66 58 115 ...
## $ Width                : num  6.62 31.75 8.3 9.9 0.3 ...
## $ Heighth              : num  16.9 19 10.2 1.3 8.9 ...
## $ Profit_margin        : num  0.15 0.25 0.08 0.08 0.09 0.05 0.05 0.05 0.05 0.05 ...
## $ Volume               : int  12 8 12 196 232 332 44 132 64 40 ...
## $ Competitors          : int  3 3 5 1 3 2 1 2 3 5 ...
## $ Professional         : int  0 0 0 0 1 0 1 1 1 1 ...
## $ Age                  : int  2 3 3 2 2 3 3 2 2 3 ...
```

Our target variable is the sales volume **Volume**. We notice that two predictors are just identifiers and can be removed: **X** and **Product_ID**. The description of the remaining 20 variables is the following:

- **Product_type**: type of electronic product (categorical).
- **Prices**: price of product (numeric).
- **X5Stars - X1Stars**: number of n-star product reviews (integer).
- **PositiveServiceReview, NegativeServiceReview**: number of positive and negative reviews of product service (integer).
- **Would_consumer_recommend_product**: score (from 0 to 1) assigned by user to the product (numeric).
- **Best_Seller_Rank**: position of product in sales ranking (integer).
- **Weight**: product weight (lbs., numeric).
- **Depth**: product depth (in., numeric).
- **Height**: product height (in., numeric).
- **Profit_margin**: profit (fraction of price, numeric).
- **Volume**: sales volume (units, integer).
- **Competitors**: number of competitor products in the market (integer).
- **Professional**: professional or business products (integer 0 or 1).
- **Age**: time of product since launch in the market (integer).

We start by simplifying the feature names:

```
# remove useless columns (index, product id),
# rename features
prod_exist %>% subset(., select= -c(X, Product_ID)) -> prod_exist
new_colnames <- c('Type', 'Price', 'x5s', 'x4s', 'x3s',
                  'x2s', 'x1s', 'PosServ', 'NegServ',
                  'Recommend', 'BestSeller', 'Weight',
                  'Depth', 'Width', 'Height', 'Profit',
                  'Vol', 'Comp', 'Prfsn', 'Age')
names(prod_exist) <- new_colnames
```

Some predictor data types need to be changed to reflect their meaning: **Professional** should be a factor with two levels (“No”, “Yes”), whereas **Depth** is a numeric variable.

```
# update predictors data types
prod_exist$Prfsn %>%
  factor(levels = c(0, 1),
         labels = c("No", "Yes")) -> prod_exist$Prfsn
```

```
prod_exist$Depth %>%
  as.character(.) %>%
  as.numeric(.) -> prod_exist$Depth
```

```
head(prod_exist)
```

```
##      Type    Price x5s x4s x3s x2s x1s PosServ NegServ Recommend
## 1      PC  949.00   3   3   2   0   0      2      0      0.9
## 2      PC 2249.99   2   1   0   0   0      1      0      0.9
## 3      PC  399.00   3   0   0   0   0      1      0      0.9
## 4  Laptop  409.99  49  19   8   3   9      7      8      0.8
## 5  Laptop 1079.99  58  31  11   7  36      7     20      0.7
## 6 Accessories 114.22  83  30  10   9  40     12      5      0.3


##  BestSeller Weight Depth Width Height Profit Vol  Comp Prfsn Age
## 1      1967   25.8 23.94  6.62  16.89   0.15  12    3    No   2
## 2      4806   50.0 35.00 31.75  19.00   0.25   8    3    No   3
## 3     12076   17.4 10.50  8.30  10.20   0.08  12    5    No   3
## 4       109    5.7 15.00  9.90   1.30   0.08 196    1    No   2
## 5       268    7.0 12.90  0.30   8.90   0.09 232    3   Yes   2
## 6        64    1.6  5.80  4.00   1.00   0.05 332    2    No   3
```

2. Cleaning and Exploration of the Dataset

We may first have a look at the distribution of the dependent variable to check for the presence of outliers that may have an outsized effect on the predictive models.


```
qqnorm(prod_exist$Vol)
```

Normal Q-Q Plot



A normalized quantile-quantile plot shows that the `Volume` target variable differs significantly from a normally distributed random variable. In particular, there are at least two points which stand out due to their huge volumes as is also apparent by looking at the boxplot below.

```
boxplot(x = prod_exist$Vol, ylab = 'sales volume')
```



These points are removed by taking the observations having sales volume < 5000 units.

```
# remove outliers
prod_exist <- filter(prod_exist, Vol < 5000)
```

Secondly, observation with NA values may also be present in the dataset.

```
# find and store NAs on separate data frame
nas <- prod_exist[!complete.cases(prod_exist), ]
dim(nas)
```

```
## [1] 16 20
```

```
summary(nas)
```

##	Type	Price	x5s	x4s
## Accessories	:6	Min. : 6.55	Min. : 0.00	Min. : 0.00
## Printer	:5	1st Qu.: 39.11	1st Qu.: 5.75	1st Qu.: 1.75
## Printer Supplies	:2	Median : 132.72	Median : 13.00	Median : 5.00
## Laptop	:1	Mean : 186.40	Mean : 38.56	Mean : 14.75
## PC	:1	3rd Qu.: 221.94	3rd Qu.: 21.00	3rd Qu.: 12.25
## Software	:1	Max. : 609.99	Max. : 349.00	Max. : 118.00
## (Other)	:0			
##	x3s	x2s	x1s	PosServ
## Min.	: 0.000	Min. : 0.000	Min. : 0.000	Min. : 0.000
## 1st Qu.	: 0.000	1st Qu.: 0.000	1st Qu.: 0.750	1st Qu.: 1.000
## Median	: 2.000	Median : 0.000	Median : 2.000	Median : 2.500
## Mean	: 4.625	Mean : 2.062	Mean : 5.875	Mean : 9.875
## 3rd Qu.	: 4.250	3rd Qu.: 2.500	3rd Qu.: 10.500	3rd Qu.: 5.000
## Max.	: 27.000	Max. : 11.000	Max. : 21.000	Max. : 64.000


```
##
##      NegServ      Recommend      BestSeller      Weight
## Min.   : 0.000   Min.   :0.500   Min.   :559   Min.   : 0.400
## 1st Qu.: 0.000   1st Qu.:0.675   1st Qu.:559   1st Qu.: 1.000
## Median : 1.000   Median :0.800   Median :559   Median : 3.805
## Mean   : 3.625   Mean   :0.750   Mean   :559   Mean   :13.632
## 3rd Qu.: 3.250   3rd Qu.:0.900   3rd Qu.:559   3rd Qu.:30.400
## Max.   :24.000   Max.   :1.000   Max.   :559   Max.   :39.000
##
##                               NA's   :15
##      Depth      Width      Height      Profit
## Min.   : 1.50   Min.   : 1.60   Min.   : 0.50   Min.   :0.0500
## 1st Qu.: 6.20   1st Qu.: 6.25   1st Qu.: 4.70   1st Qu.:0.0500
## Median :10.40   Median : 9.40   Median :11.19   Median :0.1300
## Mean   :11.26   Mean   :10.12   Mean   :10.48   Mean   :0.1412
## 3rd Qu.:16.93   3rd Qu.:14.45   3rd Qu.:14.70   3rd Qu.:0.1850
## Max.   :22.10   Max.   :20.90   Max.   :20.71   Max.   :0.3000
##
##      NA's   :1
##      Vol      Comp      Prfsn      Age
## Min.   : 0.0   Min.   :0.000   No :12   Min.   :1.000
## 1st Qu.: 30.0   1st Qu.:1.000   Yes: 4   1st Qu.:2.000
## Median : 52.0   Median :3.000           Median :3.000
## Mean   :156.6   Mean   :2.438           Mean   :2.625
## 3rd Qu.: 84.0   3rd Qu.:3.250           3rd Qu.:3.000
## Max.   :1396.0   Max.   :5.000           Max.   :4.000
##
```

There are 16 NAs, 15 of which in the `BestSeller` column, and 1 in the `Width` column. They're a relatively small number so we'll remove them from the dataset.


```
# remove NAs
prod_exist %>%
  .[complete.cases(.), ] -> prod_exist
```

Finally, to gain insight into how the observations are distributed in our dataset, let's examine the distribution of the sales volume variable against some predictors.

```
ggboxplot(data= prod_exist,
  x = 'Type',
  y = 'Vol',
  fill= 'grey') +
  theme(axis.text.x = element_text(angle = 90))
```




```
ggboxplot(data= prod_exist,
  x = 'Prfsn',
  y = 'Vol',
  fill= 'grey') +
  theme(axis.text.x = element_text(angle = 90))
```




We may expect that the sales volume be influenced by the amount of positive product reviews. The correlation between the variables volume and number of n-star reviews can be displayed using scatterplots.

```
ggscatter(data= prod_exist,
  x = 'x5s',
  y = 'Vol',
  fill= 'grey') +
labs(title = 'Sales v number of 5-star reviews')
```

```
ggscatter(data= prod_exist,  
  x = 'x4s',  
  y = 'Vol',  
  fill= 'grey') +  
labs(title = 'Sales v number of 4-star reviews')
```



The correlation is really strong for the 5-star reviews, less so for the 4-star reviews. Correlations between numerical features can be examined in more detail by calculating Pearson's correlation coefficient between feature pairs.

2.1 Quantifying and Visualizing Correlations between Variables

Determining correlations between variables is useful if we need to get rid of highly correlated predictors in order to fit the dataset with a linear model, for instance.

```
# set volume as last column in dataset
prod_exist <- prod_exist[ c(1:16, 18, 19, 20, 17) ]
```

```
# generate dummy variables for factors
dmy <- dummyVars('~ .', data = prod_exist)
prod_dmy <- prod_exist %>%
  predict(dmy, .) %>%
  data.frame()
```

```
# calculate correlations
corrData <- cor(prod_dmy)
new_colnames <-
  c('Typ.Acc', 'Typ.Disp', 'Typ.ExtW', 'Typ.GCons',
    'Typ.Lap', 'Typ.Net', 'Typ.PC', 'Typ.Prn',
    'Typ.PrnSupp', 'Typ.Smar', 'Typ.Soft',
    'Typ.Tab', 'Price', 'x5s', 'x4s',
    'x3s', 'x2s', 'x1s', 'PSer', 'NSer', 'Rec',
    'BSell', 'Wei', 'Dep', 'Wid', 'Hei', 'Prof',
```


```

'Comp', 'Prf.No', 'Prf.Yes', 'Age', 'Vol')

# shorten column and row names for plotting
colnames(corrData) <- new_colnames
rownames(corrData) <- new_colnames

# display correlogram
corrplot(corrData,
  method = 'square',
  type = 'lower',
  na.label= 'o',
  tl.srt= 30)

```



Highly correlated (collinear) predictors include:

- x5s and Vol,
- all pairs of n-star reviews predictors (x5s to x1s).

The correlation coefficients for a select pair of variables can be displayed via the following:

```
corrData["Vol", "x5s"]
```

```
## [1] 0.8342784
```

```
corrData["x5s", "x4s"]
```

```
## [1] 0.7707467
```

```
corrData["x4s", "x3s"]
```

```
## [1] 0.8740702
```

```
corrData["x3s", "x2s"]
```

```
## [1] 0.8495147
```

Some of these predictors will be removed prior to fitting a linear model to the data.

3. Model Selection and Validation: Gradient Boosted Machines